



CPSC 340 – Tutorial 1

Lironne Kurzman
lironnek@cs.ubc.ca

University of British Columbia

September 13th, 2021

GITHUB – tl;dr

Clone a repo:

```
git clone <repo url>
```

See changes:

```
git status
```

```
git diff <file name>
```

Commit changes:

```
git add <file name>
```

```
git commit -m "<commit message>" git push
```

Pull remote changes:

```
git pull
```

Notation and Gradients

- Define: \mathbb{R}^n , $x^T A x$, $\|x\|$, $\nabla f(x)$
- If $x \in \mathbb{R}^n$, $y \in \mathbb{R}^m$, $A \in \mathbb{R}^{n \times m}$ find the dimensions of:
 - $x^T A$, $x^T A y$, $x x^T$.
- If $a, x \in \mathbb{R}^n$, compute the gradients of:
 - $f(x) = a^T x$,
 - $f(x) = \log a^T x$.
- A helpful sanity check:
 - If x is a length- n (column) vector, and $f : \mathbb{R}^n \rightarrow \mathbb{R}$ then $\nabla f(x)$ is a length- n (row) vector.
- Entropy: Given a probability vector $p = [p_1, \dots, p_n]^T$,

$$\text{entropy}(p) = - \sum_{i=1}^n p_i \log(p_i).$$

Python/NumPy Basics

- numpy Arrays:

`A = np.array(...)`

- Array slicing

- Element-wise operations:

`+, -, *, /, np.amin, np.amax`

- Array shapes

- Matrix/vector operations:

`np.transpose(), @`

- Other useful things:

`np.ones, np.zeros`

Python Demo

Numpy Basics

Bonus: We can check that the gradient implementation is correct

Finite-difference approximation of the gradient

$$\epsilon = 1e-4$$

Gradient definition:

$$\frac{\partial f(x)}{\partial x_i} = \lim_{\epsilon \rightarrow 0} \frac{f(x_i + \epsilon) - f(x_i)}{\epsilon}$$

) Therefore,

$$\frac{\partial f(x)}{\partial x_i} \approx \frac{f(x_i + \epsilon) - f(x_i)}{\epsilon}$$

Bonus: Code to approximate the gradient

$$\frac{\partial f(x)}{\partial x_i} \approx \frac{f(x_i + \epsilon) - f(x_i)}{\epsilon}$$

Create a simple gradient check for single-variable function

```
def forward_diff(x, f):  
    # Approximates gradient eps = 1e-4  
    g_approx = (f(x + eps) - f(x)) / eps
```

Bonus: Code to approximate the gradient

Using scipy:

```
import numpy as np
from scipy.optimize
import approx_fprime

approx = approx_fprime(x0, foo, 1e-4)
exact = foo_grad(x0)

print("My gradient      : %s" % exact)
print("Scipy's gradient: %s" % approx)

# Assert that the two are almost equal
np.testing.assert_almost_equal(approx, exact, 3)
```